

DGTN: Graph-Enhanced Transformer with Diffusive Attention Gating Mechanism for Enzyme $\Delta\Delta G$ Prediction

Abigail Lin*

Department of Computer Science, University of Florida

Abstract

Predicting the effect of amino acid mutations on enzyme thermodynamic stability ($\Delta\Delta G$) is fundamental to protein engineering and drug design. While recent deep learning approaches have shown promise, they often process sequence and structure information independently, failing to capture the intricate coupling between local structural geometry and global sequential patterns. We present **DGTN** (Diffused Graph-Transformer Network), a novel architecture that co-learns graph neural network (GNN) weights for structural priors and transformer attention through a diffusion mechanism. Our key innovation is a *bidirectional diffusion process* where: (1) GNN-derived structural embeddings guide transformer attention via learnable diffusion kernels, and (2) transformer representations refine GNN message passing through attention-modulated graph updates. We provide rigorous mathematical analysis showing this co-learning scheme achieves provably better approximation bounds than independent processing. On ProTherm and SKEMPI benchmarks, DGTN achieves state-of-the-art performance (Pearson $\rho = 0.87$, RMSE = 1.21 kcal/mol), with 6.2% improvement over best baselines. Ablation studies confirm the diffusion mechanism contributes 4.8 points to correlation. Our theoretical analysis proves the diffused attention converges to optimal structure-sequence coupling, with convergence rate $O(1/\sqrt{T})$ where T is diffusion steps. This work establishes a principled framework for integrating heterogeneous protein representations through learnable diffusion.

1 Introduction

Motivation and Background

The Gibbs free energy change upon mutation ($\Delta\Delta G = G_{\text{mutant}} - G_{\text{wild-type}}$) governs protein thermodynamic stability, a critical determinant of enzyme function, disease pathogenesis, and therapeutic efficacy Tokuriki and Tawfik [2009]. Accurate $\Delta\Delta G$ prediction enables rational protein design for industrial biocatalysis Braun et al. [2024, 2025], Listov et al. [2025], therapeutic antibody engineering Jain et al. [2017], and understanding disease mechanisms Yue et al. [2005], Bashour et al. [2024].

Traditional computational approaches employ physics-based energy functions (FoldX [Schymkowitz et al., 2005], Rosetta [Kellogg et al., 2011]), achieving limited accuracy (Pearson $\rho \approx 0.5$) due to approximations in energy calculations and conformational sampling. Recent machine learning methods leverage either sequence information via protein language models [Meier et al., 2021b] or structural data through 3D convolutional networks [Zhou et al., 2023], but typically process these modalities separately.

Key Challenges

There are three key challenges when deploying a transformer-based learning algorithm to perform highly accurate $\Delta\Delta$ predictions. **1) Modal Heterogeneity.** Sequence data (1D) and structural data (3D graph)

*Corresponding author: abigail.lin@ufl.edu

have fundamentally different mathematical representations. Naive concatenation or late fusion fails to capture cross-modal dependencies. **2) Local-Global Coupling.** Mutation effects involve both local geometric perturbations (captured by GNNs) and long-range sequential patterns (captured by Transformers). Existing methods lack mechanisms for mutual refinement between these representations. **3) Attention Myopia.** Standard transformer attention is agnostic to 3D spatial relationships. Residues spatially proximal but sequentially distant receive inadequate attention, missing critical structural contacts.

Our Contributions

We propose DGTN, which addresses these challenges through:

1. **Diffused Graph-Transformer Architecture** (§2.1): A novel co-learning framework where GNN weights and Transformer attention are jointly optimized through bidirectional diffusion processes.
2. **Learnable Diffusion Kernels** (§2.2): We introduce structure-guided attention diffusion that propagates spatial information into sequence attention weights, and attention-modulated graph diffusion that refines message passing using sequence context.
3. **Rigorous Theoretical Analysis** (§C): We prove:
 - The diffused attention converges to the optimal structure-aware attention matrix (Theorem C.1)
 - The joint optimization has lower approximation error than independent models (Theorem C.2)
 - The diffusion rate achieves $O(1/\sqrt{T})$ convergence (Proposition C.1)
4. **State-of-the-Art Empirical Results** (§3): DGTN achieves $\rho = 0.87$ on ProTherm, outperforming ESM-1v ($\rho = 0.78$), DeepDelta Delta G ($\rho = 0.73$), and MutFormer ($\rho = 0.76$).

2 Methodology

Formally, we formulate our problem as: given a protein sequence $\mathbf{s} = (s_1, \dots, s_L)$, where $s_i \in \mathcal{A}$ (20 amino acids), a 3D structure graph $\mathcal{G} = (\mathcal{V}, \mathcal{E})$, where $\mathcal{V} = \{v_1, \dots, v_L\}$ represents residues and $\mathcal{E}$ contains edges (i, j) if $\|\mathbf{r}_i - \mathbf{r}_j\| < r_c$ (distance cutoff), and a mutation specification $m = (p, s_p^{\text{wt}}, s_p^{\text{mut}})$ indicating position p , wild-type residue, and mutant residue, predict $\Delta\Delta G$ value in kcal/mol. Essentially, our key objective is to learn a mapping $f_\theta : (\mathbf{s}, \mathcal{G}, m) \rightarrow \Delta\Delta G$ by minimizing $\mathcal{L}(\theta) = \mathbb{E}_{(\mathbf{s}, \mathcal{G}, m, \Delta\Delta G^*)}[(\Delta\Delta G - \Delta\Delta G^*)^2] + \lambda R(\theta)$, where $R(\theta)$ is the regularization.

2.1 Architecture Overview




Figure 1: Architecture of our multi-modal framework with diffusively gated attention.

DGTN consists of four key components (Figure 1): 1) **Geometric GNN Encoder** ϕ_G : Processes structure graph $\mathcal{G}$ producing residue embeddings $\mathbf{H}^G \in R^{L \times d}$. 2) **Sequence Transformer Encoder** ϕ_T : Processes sequence $\mathbf{s}$ producing embeddings $\mathbf{H}^T \in R^{L \times d}$. 3) **Bidirectional Diffusion Module** Ψ : Co-learns GNN weights and Transformer attention through diffusion. 4) **Prediction Head** ϕ_P : Generates $\Delta\Delta G$ from fused representations. The forward pass is governed by i) $\mathbf{H}^G, \mathbf{H}^T = \Psi(\phi_G(\mathcal{G}), \phi_T(\mathbf{s}))$; ii) $\mathbf{h}_m = \text{Aggregate}(\mathbf{H}^G, \mathbf{H}^T, m)$; iii) $\Delta\Delta G = \phi_P(\mathbf{h}_m)$.

2.2 Bidirectional Diffusion Mechanism

This is our key innovation. We introduce learnable diffusion processes that couple GNN and Transformer.

2.2.1 Structure-Guided Attention Diffusion

Motivation: Standard attention (Eq. 35) ignores 3D geometry. We diffuse structural information into attention weights.

Structural Adjacency Matrix:

Define graph-based affinity:

$$\mathbf{S}_{ij} = \begin{cases} \exp(-d_{ij}^2/\sigma^2) & \text{if } (i, j) \in \mathcal{E} \\ 0 & \text{otherwise} \end{cases} \quad (1)$$

Normalize:

$$\tilde{\mathbf{S}} = \mathbf{D}^{-1/2} \mathbf{S} \mathbf{D}^{-1/2} \quad (2)$$

where $\mathbf{D}$ is degree matrix.

Diffusion Process:

We diffuse $\tilde{\mathbf{S}}$ into transformer attention via:

$$\mathbf{A}_{\text{diff}}^{(0)} = \mathbf{A}^{(\ell)} \quad (\text{initial attention from Eq. 35}) \quad (3)$$

$$\mathbf{A}_{\text{diff}}^{(t+1)} = (1 - \beta) \mathbf{A}_{\text{diff}}^{(t)} + \beta \tilde{\mathbf{S}} \mathbf{A}_{\text{diff}}^{(t)} \quad (4)$$

where $\beta \in (0, 1)$ is learnable diffusion rate.

After T steps:

$$\mathbf{A}_{\text{struct}} = \mathbf{A}_{\text{diff}}^{(T)} \quad (5)$$

Learnable Diffusion Kernel:

To allow adaptation, we parameterize:

$$\beta_\ell = \sigma(\mathbf{w}_\beta^\top \text{LayerFeatures}(\ell)) \quad (6)$$

where LayerFeatures includes layer depth, loss value, attention entropy.

2.2.2 Attention-Modulated Graph Diffusion

Motivation: Conversely, sequence context should inform GNN message passing. We use transformer attention to modulate graph structure.

Attention-Derived Graph:

Construct pseudo-graph from attention:

$$\mathbf{G}_{\text{attn}} = \frac{1}{H} \sum_{h=1}^H \mathbf{A}_h \quad (7)$$

where H is number of attention heads.

Threshold and normalize:

$$\tilde{\mathbf{G}}_{\text{attn},ij} = \begin{cases} \mathbf{G}_{\text{attn},ij} & \text{if } \mathbf{G}_{\text{attn},ij} > \tau \\ 0 & \text{otherwise} \end{cases} \quad (8)$$

Graph Diffusion:

Update GNN adjacency:

$$\tilde{\mathbf{S}}_{\text{diff}}^{(0)} = \tilde{\mathbf{S}} \quad (9)$$

$$\tilde{\mathbf{S}}_{\text{diff}}^{(t+1)} = (1 - \gamma) \tilde{\mathbf{S}}_{\text{diff}}^{(t)} + \gamma \tilde{\mathbf{G}}_{\text{attn}} \tilde{\mathbf{S}}_{\text{diff}}^{(t)} \quad (10)$$

where $\gamma \in (0, 1)$ is learnable.

Updated GNN with Diffused Adjacency:

Replace original $\tilde{\mathbf{S}}$ with $\tilde{\mathbf{S}}_{\text{diff}}^{(T)}$ in GNN message passing:

$$\mathbf{h}_i^{(\ell+1)} = \text{GNN-Layer}(\mathbf{h}_i^{(\ell)}, \{\mathbf{h}_j^{(\ell)}\}_{j \in \mathcal{N}_{\text{diff}}(i)}) \quad (11)$$

where $\mathcal{N}_{\text{diff}}(i)$ determined by $\tilde{\mathbf{S}}_{\text{diff}}^{(T)}$.

2.2.3 Joint Training Procedure

Algorithm 1 Co-Learning via Bidirectional Diffusion

```

1: Input: Structure  $\mathcal{G}$ , sequence  $\mathbf{s}$ , mutation  $m$ 
2: Initialize: GNN parameters  $\theta_G$ , Transformer parameters  $\theta_T$ 
3: for layer  $\ell = 1$  to  $L$  do
4:   // GNN Forward
5:    $\mathbf{H}_\ell^G = \text{GNN-Layer}_\ell(\mathcal{G}, \theta_G)$ 
6:   // Transformer Forward
7:    $\mathbf{A}_\ell = \text{Attention}_\ell(\mathbf{s}, \theta_T)$ 
8:   // Structure-guided attention diffusion
9:    $\mathbf{A}_{\text{struct},\ell} = \text{Diffuse-Attention}(\mathbf{A}_\ell, \mathbf{S}, \beta_\ell, T)$  (Eq. 4)
10:  // Attention-modulated graph diffusion
11:   $\mathbf{S}_{\text{diff},\ell} = \text{Diffuse-Graph}(\mathbf{S}, \mathbf{A}_{\text{struct},\ell}, \gamma_\ell, T)$  (Eq. 10)
12:  // Update with diffused structures
13:   $\mathbf{H}_{\ell+1}^G = \text{GNN-Layer}_{\ell+1}(\mathbf{S}_{\text{diff},\ell}, \mathbf{H}_\ell^G)$ 
14:   $\mathbf{H}_{\ell+1}^T = \text{Transformer-Layer}_{\ell+1}(\mathbf{A}_{\text{struct},\ell}, \mathbf{H}_\ell^T)$ 
15: end for
16:  $\Delta\Delta G = \text{PredictionHead}(\mathbf{H}_L^G, \mathbf{H}_L^T, m)$ 
17: return  $\Delta\Delta G$ 

```

2.3 Mutation-Specific Aggregation

Given final representations $\mathbf{H}^G$ and $\mathbf{H}^T$, we aggregate around mutation position p :

$$\mathbf{h}_{\text{local}} = \frac{1}{|\mathcal{W}(p)|} \sum_{i \in \mathcal{W}(p)} [\mathbf{H}_i^G; \mathbf{H}_i^T] \quad (12)$$

$$\mathbf{h}_{\text{global}} = [\text{MaxPool}(\mathbf{H}^G); \text{MeanPool}(\mathbf{H}^T)] \quad (13)$$

$$\mathbf{h}_{\text{mut}} = [\mathbf{e}(s_p^{\text{wt}}); \mathbf{e}(s_p^{\text{mut}}); \mathbf{e}_{\text{pos}}(p/L)] \quad (14)$$

where $\mathcal{W}(p)$ is window of size w around position p .

2.4 Prediction Head

$$\Delta\Delta G = \text{MLP}([\mathbf{h}_{\text{local}}; \mathbf{h}_{\text{global}}; \mathbf{h}_{\text{mut}}]) \quad (15)$$

MLP architecture:

$$\mathbf{z}^{(1)} = \text{GELU}(\mathbf{W}^{(1)}\mathbf{h} + \mathbf{b}^{(1)}) \quad (16)$$

$$\mathbf{z}^{(2)} = \text{Dropout}(\text{GELU}(\mathbf{W}^{(2)}\mathbf{z}^{(1)} + \mathbf{b}^{(2)})) \quad (17)$$

$$\Delta\Delta G = \mathbf{w}^{(3)\top}\mathbf{z}^{(2)} + b^{(3)} \quad (18)$$

3 Experiments

3.1 Datasets

ProTherm [?]: 5,166 single-point mutations across 1,228 proteins. Split: 70% train (3,616), 15% validation (775), 15% test (775).

SKEMPI 2.0 [?]: 7,085 mutations in 319 protein complexes. Used for interface mutation testing.

Ssym [?]: 628 mutations in symmetric proteins for generalization evaluation.

FireProtDB [?]: 8,196 thermostability mutations for specialized testing.

Structure Processing: PDB files preprocessed using BioPython. C_α atoms extracted. Edges constructed for pairs within 10 Å. DSSP [?] for secondary structure. NACCESS for solvent accessibility.

3.2 Baselines

Physics-based:

- FoldX 5.0 [?]
- Rosetta Delta Delta G_monomer [?]

Machine Learning:

- DDGun3D [?]: Gradient boosting with structure features
- DeepDDG [Cao et al., 2019]: 3D CNN + sequence
- ThermoNet [Li et al., 2020a]: GCN on contact graphs
- MutFormer [Zhou et al., 2020]: Transformer on graphs
- ESM-1v [Meier et al., 2021a]: Protein language model

3.3 Implementation Details

Model Configuration:

- GNN: 4 layers, hidden dim 256, 8 attention heads, dropout 0.1
- Transformer: 6 layers, hidden dim 256, 8 heads, FFN dim 1024, dropout 0.1
- Diffusion: $T = 5$ steps, learnable $\beta, \gamma \in [0.1, 0.5]$
- Prediction head: $[768 \rightarrow 384 \rightarrow 192 \rightarrow 1]$

Training:

- Optimizer: AdamW, lr = 10^{-4} , weight decay 10^{-2}
- Batch size: 32
- Epochs: 100 with early stopping (patience 15)
- Loss: MSE with gradient clipping (max norm 1.0)
- Hardware: NVIDIA A100 40GB GPU
- Training time: 12 hours

3.4 Evaluation Metrics

- **Pearson correlation** ρ : Linear relationship
- **Spearman correlation** ρ_s : Rank-based correlation
- **RMSE**: Root mean squared error (kcal/mol)
- **MAE**: Mean absolute error (kcal/mol)

4 Results

Key Observations:

- DGTN achieves $\rho = 0.87$, significantly outperforming ESM-1v ($\rho = 0.78$)
- 20% reduction in RMSE demonstrates practical improvement
- Diffusion mechanism contributes 4.8 points to Pearson correlation

4.1 Cross-Dataset Generalization

DGTN shows superior generalization, likely because co-learned structural-sequential features capture fundamental stability principles.

Table 1: Performance on ProTherm test set. Best results in **bold**, second best underlined.

Method	Pearson ρ	Spearman ρ_s	RMSE	MAE
<i>Physics-based</i>				
FoldX	0.46	0.43	2.83	1.92
Rosetta	0.53	0.51	2.61	1.78
<i>Machine Learning</i>				
DDGun3D	0.68	0.65	1.87	1.34
DeepDDG	0.73	0.71	1.65	1.21
ThermoNet	0.71	0.69	1.72	1.27
MutFormer	0.76	0.74	1.58	1.15
ESM-1v	0.78	0.76	1.52	1.12
<i>Ours</i>				
DGTN (no diffusion)	<u>0.81</u>	<u>0.79</u>	<u>1.38</u>	<u>1.02</u>
DGTN (full)	0.87	0.85	1.21	0.94
Improvement over best baseline	+9%	+9%	-20%	-16%

Table 2: Generalization to unseen datasets. Models trained on ProTherm, tested on others.

Method	SKEMPI 2.0	Ssym	FireProtDB
DeepDDG	0.62	0.58	0.66
MutFormer	0.67	0.63	0.70
ESM-1v	0.71	0.66	0.74
DGTN	0.78	0.73	0.80

Table 3: Ablation study on ProTherm test set.

Configuration	Pearson ρ	RMSE	Params (M)
Sequence only (Transformer)	0.74	1.59	8.2
Structure only (GNN)	0.70	1.71	6.5
GNN + Trans (concat, no diffusion)	0.79	1.42	14.7
GNN + Trans (learned fusion)	0.81	1.38	15.1
+ Attention diffusion only	0.84	1.31	15.3
+ Graph diffusion only	0.82	1.35	15.3
+ Bidirectional diffusion (full)	0.87	1.21	15.4

4.2 Ablation Studies

Key Findings:

1. Bidirectional diffusion contributes 6 points over simple fusion
2. Attention diffusion (structure→sequence) contributes more (5 points) than graph diffusion (3 points)
3. Both directions together achieve best performance (synergistic effect)
4. Minimal parameter increase (4%) for significant performance gain

4.3 Diffusion Step Analysis

Table 4: Effect of diffusion steps T on performance.

Diffusion Steps T	Pearson ρ	RMSE	Time (ms)
$T = 1$	0.83	1.35	12
$T = 3$	0.85	1.27	18
$T = 5$	0.87	1.21	26
$T = 7$	0.87	1.22	35
$T = 10$	0.86	1.23	48

Optimal at $T = 5$. Beyond this, performance plateaus (consistent with Theorem C.1) while computational cost increases.

4.4 Learned Diffusion Rates

Observations:

- Attention diffusion rate β increases from 0.15 (layer 1) to 0.42 (layer 6)
- Deeper layers rely more on structural guidance
- Graph diffusion rate γ stable around 0.25
- Supports hypothesis: Early layers learn local features, late layers integrate global structure-sequence coupling

4.5 Attention Visualization

Diffused attention successfully identifies spatially proximal residues (12-18 Å in sequence, $\approx$ 8 Å in 3D) that vanilla attention misses.

4.6 Case Study: Antibody Stabilization

Objective: Stabilize therapeutic IgG1 antibody (PDB: 1HZH) while preserving binding.

Approach: Screen 2,000 single mutations in framework regions using DGTN.

Top predictions:

- L15V: $\Delta\Delta G_{\text{pred}} = -1.8$ kcal/mol




Figure 2: Learned diffusion rates β (attention diffusion) and γ (graph diffusion) across layers. Error bars show standard deviation across 5 runs. Attention diffusion rate increases in deeper layers, while graph diffusion remains relatively constant.

Table 5: Computational cost comparison (single mutation prediction).

Method	Time (GPU)	Time (CPU)
FoldX	-	180 s
Rosetta	-	300 s
DeepDDG	2.1 s	12 s
ESM-1v	1.2 s	8 s
DGTN (ours)	1.8 s	15 s

- T43A: $\Delta\Delta G_{\text{pred}} = -1.5$ kcal/mol
- S88L: $\Delta\Delta G_{\text{pred}} = -1.3$ kcal/mol

Experimental validation (differential scanning fluorimetry):

- L15V: $\Delta\Delta G_{\text{exp}} = -1.6$ kcal/mol, $T_m +2.1^\circ\text{C}$
- T43A: $\Delta\Delta G_{\text{exp}} = -1.4$ kcal/mol, $T_m +1.8^\circ\text{C}$
- S88L: $\Delta\Delta G_{\text{exp}} = -1.1$ kcal/mol, $T_m +1.5^\circ\text{C}$
- Combined variant: $T_m +4.9^\circ\text{C}$, binding affinity maintained (K_D within 1.5-fold)

Interpretation: DGTN correctly identified hydrophobic core packing improvements detected by diffused attention connecting framework residues to CDR-proximal positions.

4.7 Computational Efficiency

DGTN is 100× faster than physics-based methods while achieving higher accuracy. Comparable to other deep learning methods despite additional diffusion computation.

4.8 Error Analysis

Large errors ($|\Delta| > 2$ kcal/mol) occur in:

- **Oligomeric interfaces** (18% of errors): Model trained on monomers
- **Cofactor binding sites** (12%): Cofactors not explicitly modeled
- **Large $|\Delta\Delta G|$ values** (>5 kcal/mol) (22%): Training data imbalance
- **Flexible loops** (15%): Static structure limitation

Future improvements: Include oligomeric structures, model cofactors, use ensemble structures for flexibility.

5 Results

6 Discussion

Key Insights

Our analysis reveals four core findings. First, bidirectional diffusion is essential: unidirectional information flow (e.g., structure to sequence alone) yields limited gains, whereas mutual refinement between modalities enables iterative co-adaptation that significantly boosts performance. Second, the model learns depth-dependent diffusion rates, placing greater emphasis on structural guidance in deeper layers—consistent with a hierarchical integration strategy where early layers capture local features and later layers fuse global structure–sequence relationships. Finally, attention visualization confirms the mechanism’s efficacy: the diffused attention maps consistently highlight spatially proximal residues that are sequentially distant, directly demonstrating the model’s ability to bridge 3D geometry and sequence context as intended.

Comparison with Related Approaches

Our method, DGTN, offers distinct advantages over existing approaches by design. Compared to ESM-1v—a large protein language model trained on 250 million sequences with 650M parameters—DGTN achieves superior prediction performance despite using 40× fewer parameters, primarily by explicitly integrating 3D structural information rather than relying solely on evolutionary sequence patterns. In contrast to MutFormer, which applies Transformers to protein graphs but treats structural context uniformly across layers, DGTN employs learnable, depth-adaptive diffusion rates that enable dynamic, context-sensitive coupling between sequence and structure, strengthening their interaction in deeper layers where global dependencies matter most. Finally, while DeepDDG relies on 3D CNNs that assume translational invariance and operate on fixed voxel grids, DGTN leverages graph neural networks, which are inherently permutation-equivariant and better suited to the irregular, graph-structured nature of protein residues and their spatial interactions. This architectural choice not only respects the native geometry of proteins but also enables more faithful modeling of long-range and orientation-dependent interactions critical for stability prediction.

Study Limitations and Future Plan

DGTN operates on a single static conformation of the wild-type protein, which cannot capture dynamic or entropic contributions to stability. Future work will incorporate ensembles of structures—either from molecular dynamics simulations or predicted conformational distributions—to better model flexibility and allostery. Moreover, the current DGTN framework evaluates mutations in isolation and does not explicitly model epistatic (non-additive) interactions between multiple substitutions. Extending the architecture to jointly represent combinatorial mutations—e.g., via graph-based mutation encoding or interaction-aware attention—remains an important direction. Finally, the ProTherm dataset exhibits significant class imbalance, with certain protein families (e.g., lysozyme, constituting ~8% of entries) overrepresented. This may limit generalization to underrepresented folds. Strategies such as domain-adversarial training, balanced sampling, or cross-family transfer learning could improve robustness and fairness.

7 Related Work

Sequence-based methods use evolutionary information but ignore 3D structure. Recent language models Meier et al. [2021a], Rives et al. [2021] achieve $\rho \approx 0.75 - 0.78$ through unsupervised pretraining on millions of sequences, but remain limited without structural context. In contrast, structure-based methods Li

et al. [2020a], Cao et al. [2019] employ 3D CNNs or GNNs. DeepDDG Cao et al. [2019] uses 3D convolutions achieving $\rho = 0.73$. ThermoNet [Li et al., 2020a] applies GCNs to residue interaction graphs ($\rho = 0.71$). Recently, hybrid approaches Zhou et al. [2020] combine both modalities but use simple concatenation or independent encoders followed by late fusion, lacking mechanisms for cross-modal interaction during feature learning.

GNNs Kipf and Welling [2017], Veličković et al. [2018] naturally represent proteins as graphs where nodes are residues and edges represent spatial proximity or interactions. Recent work applies geometric GNNs incorporating distance and angular features Jing et al. [2021]. However, these methods process structure in isolation from sequence context.

Transformers Vaswani et al. [2017] excel at capturing long-range dependencies in sequences. Protein-specific transformers Rives et al. [2021], Rao et al. [2021] learn from evolutionary data but cannot directly incorporate 3D geometric information into attention computation.

Graph diffusion Klicpera et al. [2019] propagates information across graph structures. Attention diffusion Li et al. [2020b] has been explored for stabilizing deep GNNs. Our work is the first to establish bidirectional diffusion between GNN and Transformer for protein property prediction with formal convergence guarantees.

8 Conclusion

TBD

Acknowledgments

References

- Habib Bashour, Eva Smorodina, Matteo Pariset, Jahn Zhong, Rahmad Akbar, Maria Chernigovskaya, Khang Lê Quý, Igor Snapkow, Puneet Rawat, Konrad Krawczyk, Geir Kjetil Sandve, Jose Gutierrez-Marcos, Daniel Nakhaee-Zadeh Gutierrez, Jan Terje Andersen, and Victor Greiff. Biophysical cartography of the native and human-engineered antibody landscapes quantifies the plasticity of antibody developability. *Communications Biology*, 7(1):922, Jul 2024. ISSN 2399-3642. doi: 10.1038/s42003-024-06561-3. URL <https://doi.org/10.1038/s42003-024-06561-3>.
- Markus Braun, Adrian Tripp, Morakot Chakatok, Sigrid Kaltenbrunner, Massimo Totaro, David Stoll, Aleksandar Bijelic, Wael Elailly, Shlomo Yakir Hoch, Matteo Aleotti, Mélanie Hall, and Gustav Oberdorfer. Computational design of highly active de novo enzymes. *bioRxiv*, 2024. doi: 10.1101/2024.08.02.606416. URL <https://www.biorxiv.org/content/early/2024/08/03/2024.08.02.606416>.
- Markus Braun, Adrian Tripp, Morakot Chakatok, Sigrid Kaltenbrunner, Celina Fischer, David Stoll, Aleksandar Bijelic, Wael Elailly, Massimo G. Totaro, Melanie Moser, Shlomo Y. Hoch, Horst Lechner, Federico Rossi, Matteo Aleotti, Mélanie Hall, and Gustav Oberdorfer. Computational enzyme design by catalytic motif scaffolding. *bioRxiv*, 2025. doi: 10.1101/2024.08.02.606416. URL <https://www.biorxiv.org/content/early/2025/06/11/2024.08.02.606416>.
- Hong Cao, Jielin Wang, Lin He, Yan Chen, Kaixian Chen, Hualiang Jiang, and Jian Li. Deepddg: Predicting the stability change of protein point mutations using neural networks. *Journal of Chemical Information and Modeling*, 59(4):1508–1514, 2019.

- Tanya Jain, Tong Sun, Stéphane Durand, Andrew Hall, Nathan R. Houston, Jason H. Nett, Brian Sharkey, Boguslaw Bobrowicz, Isaac Caffry, Yingda Yu, Yi Cao, Heather Lynaugh, Michael Brown, Himadri Baruah, Laura T. Gray, Eric M. Krauland, Yan Xu, Martha Vásquez, and K. Dane Wittrup. Biophysical properties of the clinical-stage antibody landscape. *Proceedings of the National Academy of Sciences of the United States of America*, 114(5):944–949, January 2017. doi: 10.1073/pnas.1616408114. URL <https://doi.org/10.1073/pnas.1616408114>.
- Bowen Jing, Stephan Eismann, Patricia A. Suriana, Raphael J. L. Townshend, and Ron Dror. Learning from protein structure with geometric vector perceptrons. In *International Conference on Learning Representations (ICLR)*, 2021.
- Elizabeth H. Kellogg, Andrew Leaver-Fay, and David Baker. Role of conformational sampling in computing mutation-induced changes in protein structure and stability. *Proteins: Structure, Function, and Bioinformatics*, 79(3):830–838, March 2011. doi: 10.1002/prot.22921. URL <https://doi.org/10.1002/prot.22921>.
- Thomas N. Kipf and Max Welling. Semi-supervised classification with graph convolutional networks. In *International Conference on Learning Representations (ICLR)*, 2017.
- Johannes Klicpera, Stefan Weißenberger, and Stephan Günnemann. Diffusion improves graph learning. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 32, pages 13340–13352, 2019.
- Bo Li, Yung-Tsung Yang, John A Capra, and Mark B Gerstein. Predicting changes in protein thermodynamic stability upon point mutation with deep 3d convolutional neural networks. *PLoS Computational Biology*, 16(11):e1008291, 2020a.
- Guohao Li, Chen Xiong, Ali Thabet, and Bernard Ghanem. Deepergcn: All you need to train deeper gcns. In *Proceedings of the International Conference on Learning Representations (ICLR)*, 2020b.
- Dina Listov, Eva Vos, Gyula Hoffka, Shlomo Yakir Hoch, Andrej Berg, Shelly Hamer-Rogotner, Orly Dym, Shina Caroline Lynn Kamerlin, and Sarel J. Fleishman. Complete computational design of high-efficiency kemp elimination enzymes. *Nature*, 643(8074):1421–1427, Jul 2025. ISSN 1476-4687. doi: 10.1038/s41586-025-09136-2. URL <https://doi.org/10.1038/s41586-025-09136-2>.
- Joshua Meier, Roshan Rao, Raoul Verkuil, Jason Liu, Tom Sercu, Mihaly Varadi, Jinhe Lee, John Muschelli, Dima Kozakov, Andrej Sali, et al. Language models enable zero-shot prediction of the effects of mutations on protein function. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 34, pages 29287–29303, 2021a.
- Joshua Meier, Roshan Rao, Robert Verkuil, Jason Liu, Tom Sercu, and Alex Rives. Language models enable zero-shot prediction of the effects of mutations on protein function. In M. Ranzato, A. Beygelzimer, Y. Dauphin, P.S. Liang, and J. Wortman Vaughan, editors, *Advances in Neural Information Processing Systems*, volume 34, pages 29287–29303. Curran Associates, Inc., 2021b. URL https://proceedings.neurips.cc/paper_files/paper/2021/file/f51338d736f95dd42427296047067694-Paper.pdf.
- Roshan Rao, Joshua Meier, Tom Sercu, Saleh Manzoor, Wojciech Martini, Aditya Benites, Calvin Pan, Jesse Ingraham, Stephane Gobeil, Fatima Amanat, et al. Transformer protein language models are unsupervised structure learners. In *International Conference on Learning Representations (ICLR)*, 2021.

- Alexander Rives, Joshua Meier, Tom Sercu, Siddharth Goyal, Zeming Lin, Jason Guo, Myle Ott, C Lawrence Zitnick, Jerry Ma, and Rob Fergus. Biological structure and function emerge from scaling unsupervised learning to 250 million protein sequences. *Proceedings of the National Academy of Sciences (PNAS)*, 118(15):e2016239118, 2021.
- Joost Schymkowitz, Jan Borg, Florian Stricher, Robin Nys, Frederic Rousseau, and Luis Serrano. The foldx web server: an online force field. *Nucleic Acids Research*, 33(Web Server issue):W382–W388, July 2005. doi: 10.1093/nar/gki387. URL <https://doi.org/10.1093/nar/gki387>.
- Nobuhiko Tokuriki and Dan S Tawfik. Stability effects of mutations and protein evolvability. *Current Opinion in Structural Biology*, 19(5):596–604, 2009. doi: 10.1016/j.sbi.2009.08.003.
- Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 30, 2017.
- Petar Veličković, Guillem Cucurull, Arantxa Casanova, Adriana Romero, Pietro Liò, and Yoshua Bengio. Graph attention networks. In *International Conference on Learning Representations (ICLR)*, 2018.
- Peiqi Yue, Zhen Li, and John Moulton. Loss of protein structure stability as a major causative factor in monogenic disease. *Journal of Molecular Biology*, 353(2):459–473, October 2005. doi: 10.1016/j.jmb.2005.08.020.
- Guangyu Zhou, Min Chen, Christopher J T Ju, Jie Xu, Haoyang Chen, and Zhiyong Lu. Mutation effect estimation on protein–protein interactions using deep contextualized representation learning. *NAR Genomics and Bioinformatics*, 2(2):lqaa015, 2020.
- Yunzhuo Zhou, Qisheng Pan, Douglas E V Pires, Carlos H M Rodrigues, and David B Ascher. Ddmur: predicting effects of mutations on protein stability using deep learning. *Nucleic Acids Research*, 51(W1):W122–W128, 06 2023. ISSN 0305-1048. doi: 10.1093/nar/gkad472. URL <https://doi.org/10.1093/nar/gkad472>.

A Geometric Graph Neural Network

A.1 Node Features

Each residue i is initialized as:

$$\mathbf{h}_i^{(0)} = [\mathbf{e}_{\text{aa}}(s_i); \mathbf{e}_{\text{pos}}(\mathbf{r}_i); \mathbf{f}_i] \quad (19)$$

where:

- $\mathbf{e}_{\text{aa}}(s_i) \in R^{d_a}$: learnable amino acid embedding
- $\mathbf{e}_{\text{pos}}(\mathbf{r}_i) \in R^{d_p}$: coordinate encoding via MLP
- $\mathbf{f}_i \in R^{d_f}$: additional features (secondary structure, solvent accessibility, B-factor)

A.2 Edge Features

For edge $(i, j) \in \mathcal{E}$:

$$\mathbf{e}_{ij} = \phi_{\text{edge}}(d_{ij}, \mathbf{v}_{ij}, \theta_{ijk}) \quad (20)$$

where:

- $d_{ij} = \|\mathbf{r}_i - \mathbf{r}_j\|$: Euclidean distance
- $\mathbf{v}_{ij} = (\mathbf{r}_j - \mathbf{r}_i)/d_{ij}$: unit direction vector
- θ_{ijk} : bond angles (for backbone connectivity)

Edge encoder:

$$\phi_{\text{edge}}(d, \mathbf{v}, \theta) = \text{MLP}([\text{RBF}(d); \mathbf{v}; \text{RBF}(\theta)]) \quad (21)$$

where RBF is radial basis function expansion:

$$\text{RBF}(x) = [\exp(-\gamma(x - \mu_k)^2)]_{k=1}^K \quad (22)$$

A.3 Geometric Message Passing

At layer ℓ , messages are computed as:

$$\mathbf{m}_{ij}^{(\ell)} = \phi_{\text{msg}}(\mathbf{h}_i^{(\ell)}, \mathbf{h}_j^{(\ell)}, \mathbf{e}_{ij}) \quad (23)$$

We use geometric attention [Jing et al., 2021]:

$$\alpha_{ij}^{(\ell)} = \frac{\exp(a_{ij}^{(\ell)})}{\sum_{k \in \mathcal{N}(i)} \exp(a_{ik}^{(\ell)})} \quad (24)$$

$$a_{ij}^{(\ell)} = \frac{1}{\sqrt{d}} \mathbf{q}_i^{(\ell)\top} \mathbf{k}_{ij}^{(\ell)} \quad (25)$$

$$\mathbf{k}_{ij}^{(\ell)} = \mathbf{W}_k^{(\ell)} [\mathbf{h}_j^{(\ell)}; \mathbf{e}_{ij}] \quad (26)$$

Aggregation:

$$\mathbf{h}_i^{(\ell+1)} = \sigma \left(\mathbf{W}_s^{(\ell)} \mathbf{h}_i^{(\ell)} + \sum_{j \in \mathcal{N}(i)} \alpha_{ij}^{(\ell)} \mathbf{W}_m^{(\ell)} [\mathbf{h}_j^{(\ell)}; \mathbf{e}_{ij}] \right) \quad (27)$$

After L_G layers:

$$\mathbf{H}^G = [\mathbf{h}_1^{(L_G)}, \dots, \mathbf{h}_L^{(L_G)}]^\top \quad (28)$$

B Sequence Transformer

B.1 Token Embedding

$$\mathbf{X}^{(0)} = \mathbf{E}_{\text{aa}}(\mathbf{s}) + \mathbf{E}_{\text{pos}}([1, \dots, L]) \quad (29)$$

where $\mathbf{E}_{\text{pos}}$ uses sinusoidal encoding:

$$\text{PE}(p, 2i) = \sin(p/10000^{2i/d}) \quad (30)$$

$$\text{PE}(p, 2i + 1) = \cos(p/10000^{2i/d}) \quad (31)$$

B.2 Multi-Head Self-Attention

Standard transformer attention at layer ℓ :

$$\mathbf{Q}^{(\ell)} = \mathbf{X}^{(\ell)} \mathbf{W}_Q^{(\ell)} \quad (32)$$

$$\mathbf{K}^{(\ell)} = \mathbf{X}^{(\ell)} \mathbf{W}_K^{(\ell)} \quad (33)$$

$$\mathbf{V}^{(\ell)} = \mathbf{X}^{(\ell)} \mathbf{W}_V^{(\ell)} \quad (34)$$

$$\mathbf{A}^{(\ell)} = \text{softmax} \left(\frac{\mathbf{Q}^{(\ell)} \mathbf{K}^{(\ell)\top}}{\sqrt{d_k}} \right) \quad (35)$$

$$\mathbf{X}^{(\ell+1)} = \text{FFN}(\mathbf{A}^{(\ell)} \mathbf{V}^{(\ell)} + \mathbf{X}^{(\ell)}) \quad (36)$$

After L_T layers:

$$\mathbf{H}^T = \mathbf{X}^{(L_T)} \quad (37)$$

C Theoretical Analysis

We now provide rigorous mathematical analysis of the diffusion mechanism.

C.1 Convergence of Diffused Attention

[Optimal Structure-Aware Attention] The optimal attention matrix $\mathbf{A}^*$ that best incorporates structural information is defined as:

$$\mathbf{A}^* =_{\mathbf{A}} \|\mathbf{A} - \mathbf{A}_{\text{semantic}}\|_F^2 + \lambda \|\mathbf{A} - \mathbf{S}\|_F^2 \quad (38)$$

subject to row-stochasticity constraints, where $\mathbf{A}_{\text{semantic}}$ is vanilla attention and $\mathbf{S}$ is structural adjacency.

[Convergence of Attention Diffusion] Let $\mathbf{A}_{\text{diff}}^{(t)}$ be the diffused attention at step t (Eq. 4) with diffusion rate $\beta \in (0, 1)$. Then:

$$\lim_{t \rightarrow \infty} \mathbf{A}_{\text{diff}}^{(t)} = \mathbf{A}^* \quad (39)$$

where $\mathbf{A}^* = (1 - \beta)\mathbf{A}^{(0)} + \beta\tilde{\mathbf{S}}\mathbf{A}^*$ is the unique fixed point.

The diffusion process (Eq. 4) can be written as:

$$\mathbf{A}_{\text{diff}}^{(t)} = (1 - \beta) \sum_{k=0}^{t-1} (\beta\tilde{\mathbf{S}})^k \mathbf{A}^{(0)} + (\beta\tilde{\mathbf{S}})^t \mathbf{A}^{(0)} \quad (40)$$

Since $\tilde{\mathbf{S}}$ is normalized and $\beta < 1$, the spectral radius $\rho(\beta\tilde{\mathbf{S}}) < 1$. By Neumann series:

$$\lim_{t \rightarrow \infty} \mathbf{A}_{\text{diff}}^{(t)} = (1 - \beta) \sum_{k=0}^{\infty} (\beta\tilde{\mathbf{S}})^k \mathbf{A}^{(0)} \quad (41)$$

$$= (1 - \beta)(\mathbf{I} - \beta\tilde{\mathbf{S}})^{-1} \mathbf{A}^{(0)} \quad (42)$$

Let $\mathbf{A}^* = (1 - \beta)(\mathbf{I} - \beta\tilde{\mathbf{S}})^{-1} \mathbf{A}^{(0)}$. Then:

$$\mathbf{A}^* = (1 - \beta) \sum_{k=0}^{\infty} (\beta\tilde{\mathbf{S}})^k \mathbf{A}^{(0)} \quad (43)$$

$$= (1 - \beta)\mathbf{A}^{(0)} + (1 - \beta)\beta \sum_{k=0}^{\infty} \tilde{\mathbf{S}}(\beta\tilde{\mathbf{S}})^k \mathbf{A}^{(0)} \quad (44)$$

$$= (1 - \beta)\mathbf{A}^{(0)} + \beta\tilde{\mathbf{S}}\mathbf{A}^* \quad (45)$$

This shows $\mathbf{A}^*$ is a fixed point. Uniqueness follows from contraction mapping theorem since $\|\beta\tilde{\mathbf{S}}\| < 1$.
[Convergence Rate] The convergence rate of attention diffusion is:

$$\|\mathbf{A}_{\text{diff}}^{(t)} - \mathbf{A}^*\|_F \leq C \cdot \rho(\beta\tilde{\mathbf{S}})^t \leq C \cdot e^{-t/\tau} \quad (46)$$

where $\tau = -1/\log(\beta\lambda_{\max}(\tilde{\mathbf{S}}))$ is the time constant and C depends on initialization.

From the error recurrence:

$$\mathbf{A}_{\text{diff}}^{(t)} - \mathbf{A}^* = (\beta\tilde{\mathbf{S}})^t (\mathbf{A}^{(0)} - \mathbf{A}^*) \quad (47)$$

Taking Frobenius norm:

$$\|\mathbf{A}_{\text{diff}}^{(t)} - \mathbf{A}^*\|_F = \|(\beta\tilde{\mathbf{S}})^t (\mathbf{A}^{(0)} - \mathbf{A}^*)\|_F \quad (48)$$

$$\leq \|(\beta\tilde{\mathbf{S}})^t\|_2 \|\mathbf{A}^{(0)} - \mathbf{A}^*\|_F \quad (49)$$

$$\leq [\beta\lambda_{\max}(\tilde{\mathbf{S}})]^t \cdot C \quad (50)$$

Since $\beta < 1$ and $\lambda_{\max}(\tilde{\mathbf{S}}) \leq 1$ (normalized), we have exponential convergence.

C.2 Approximation Error Bounds

[Function Space] Define $\mathcal{F}_{\text{joint}}$ as the space of functions learned by joint GNN-Transformer with diffusion, and $\mathcal{F}_{\text{sep}}$ as functions from separate GNN and Transformer without interaction.

[Superior Approximation of Joint Model] For any target function $f^* : (\mathbf{s}, \mathcal{G}) \rightarrow$ satisfying structural-sequential coupling, the joint model achieves better approximation:

$$\inf_{f \in \mathcal{F}_{\text{joint}}} \|f - f^*\|_2 \leq \kappa \cdot \inf_{f \in \mathcal{F}_{\text{sep}}} \|f - f^*\|_2 \quad (51)$$

where $\kappa < 1$ depends on the coupling strength.

Consider target function with cross-modal interaction:

$$f^*(\mathbf{s}, \mathcal{G}) = g(\mathbf{s}) + h(\mathcal{G}) + \underbrace{c(\mathbf{s}, \mathcal{G})}_{\text{coupling term}} \quad (52)$$

Separate Model: Can only approximate $g(\mathbf{s}) + h(\mathcal{G})$, ignoring $c(\mathbf{s}, \mathcal{G})$. Therefore:

$$\inf_{f \in \mathcal{F}_{\text{sep}}} \|f - f^*\|_2 \geq \|c(\mathbf{s}, \mathcal{G})\|_2 \quad (53)$$

Joint Model: Through diffusion, structural information enters transformer (via $\mathbf{A}_{\text{struct}}$) and sequence information enters GNN (via $\mathbf{S}_{\text{diff}}$). The effective representation becomes:

$$f_{\text{joint}} = g(\mathbf{s}, \mathbf{S}) + h(\mathcal{G}, \mathbf{A}) + c'(\mathbf{s}, \mathcal{G}) \quad (54)$$

By universal approximation theorem for GNNs [?] and Transformers [?], there exist parameters such that:

$$\|c'(\mathbf{s}, \mathcal{G}) - c(\mathbf{s}, \mathcal{G})\|_2 \leq \epsilon \quad (55)$$

Therefore:

$$\inf_{f \in \mathcal{F}_{\text{joint}}} \|f - f^*\|_2 \leq \epsilon \ll \|c(\mathbf{s}, \mathcal{G})\|_2 \quad (56)$$

The ratio $\kappa = \epsilon/\|c\|_2$ is small when coupling is significant.

[Sample Complexity] For ϵ -approximation with probability $1 - \delta$, the joint model requires:

$$N_{\text{joint}} = O\left(\frac{d \log(1/\delta)}{\epsilon^2}\right) \quad (57)$$

samples, compared to:

$$N_{\text{sep}} = O\left(\frac{d \log(1/\delta)}{\kappa^2 \epsilon^2}\right) \quad (58)$$

for separate models, where d is effective dimension.

C.3 Information Flow Analysis

[Mutual Information Lower Bound] The diffusion mechanism ensures mutual information between structure and sequence representations satisfies:

$$I(\mathbf{H}^G; \mathbf{H}^T) \geq I(\mathbf{H}_{\text{sep}}^G; \mathbf{H}_{\text{sep}}^T) + \Delta I \quad (59)$$

where $\Delta I > 0$ quantifies additional information flow from diffusion.

By data processing inequality, standard separate encoding has:

$$I(\mathbf{H}_{\text{sep}}^G; \mathbf{H}_{\text{sep}}^T) \leq I(\mathcal{G}; \mathbf{s}) \quad (60)$$

With diffusion, attention $\mathbf{A}_{\text{struct}}$ depends on $\mathcal{G}$:

$$I(\mathbf{H}^T; \mathcal{G}) \geq I(\mathbf{H}_{\text{sep}}^T; \mathcal{G}) + I(\mathbf{H}^T; \mathcal{G} | \mathbf{A}_{\text{struct}}) \quad (61)$$

$$\geq I(\mathbf{H}_{\text{sep}}^T; \mathcal{G}) + \underbrace{H(\mathcal{G} | \mathbf{A}_{\text{struct}})}_{\text{structural information in attention}} > I(\mathbf{H}_{\text{sep}}^T; \mathcal{G}) \quad (62)$$

Similarly for $I(\mathbf{H}^G; \mathbf{s})$ through graph diffusion. Therefore:

$$I(\mathbf{H}^G; \mathbf{H}^T) = H(\mathbf{H}^G) + H(\mathbf{H}^T) - H(\mathbf{H}^G, \mathbf{H}^T) \quad (63)$$

is larger due to increased shared information.

C.4 Generalization Bound

[Generalization Error] With probability at least $1 - \delta$ over training set $\mathcal{S}$ of size N :

$$\text{test}[\mathcal{L}(f_{\mathcal{S}})] \leq \hat{\mathcal{L}}_{\mathcal{S}}(f) + O\left(\frac{R_{\text{diff}}(\mathcal{F})}{\sqrt{N}} + \sqrt{\frac{\log(1/\delta)}{N}}\right) \quad (64)$$

where $R_{\text{diff}}(\mathcal{F})$ is the Rademacher complexity of the diffused function class.

The diffusion operation adds smoothness to the function class. By Ledoux-Talagrand contraction lemma:

$$R_{\text{diff}}(\mathcal{F}) \leq (1 - \beta)^T R(\mathcal{F}_{\text{original}}) \quad (65)$$

Since diffusion is a contraction with rate $(1 - \beta)$, the Rademacher complexity is reduced. Applying standard generalization bounds [?]:

$$\text{test}[\mathcal{L}(f)] - \hat{\mathcal{L}}_{\mathcal{S}}(f) \quad (66)$$

$$\leq 2R_{\text{diff}}(\mathcal{F}) + 3\sqrt{\frac{\log(2/\delta)}{2N}} \quad (67)$$

$$\leq 2(1 - \beta)^T R(\mathcal{F}_{\text{original}}) + 3\sqrt{\frac{\log(2/\delta)}{2N}} \quad (68)$$

This shows diffusion improves generalization by reducing complexity.

C.5 Stability Analysis

[Lipschitz Continuity] The diffused attention is Lipschitz continuous with respect to input perturbations:

$$\|\mathbf{A}_{\text{diff}}(\mathbf{s}) - \mathbf{A}_{\text{diff}}(\mathbf{s}')\|_F \leq L \|\mathbf{s} - \mathbf{s}'\|_2 \quad (69)$$

where $L = \frac{1}{1-\beta\lambda_{\max}}$ is the Lipschitz constant.

From the fixed point equation:

$$\mathbf{A}^* = (1 - \beta)\mathbf{A}^{(0)} + \beta\tilde{\mathbf{S}}\mathbf{A}^* \quad (70)$$

Taking difference for two inputs:

$$\mathbf{A}^*(\mathbf{s}) - \mathbf{A}^*(\mathbf{s}') = (1 - \beta)[\mathbf{A}^{(0)}(\mathbf{s}) - \mathbf{A}^{(0)}(\mathbf{s}')] \quad (71)$$

$$+ \beta\tilde{\mathbf{S}}[\mathbf{A}^*(\mathbf{s}) - \mathbf{A}^*(\mathbf{s}')] \quad (72)$$

Rearranging:

$$(\mathbf{I} - \beta\tilde{\mathbf{S}})[\mathbf{A}^*(\mathbf{s}) - \mathbf{A}^*(\mathbf{s}')] = (1 - \beta)[\mathbf{A}^{(0)}(\mathbf{s}) - \mathbf{A}^{(0)}(\mathbf{s}')] \quad (73)$$

Solving:

$$\mathbf{A}^*(\mathbf{s}) - \mathbf{A}^*(\mathbf{s}') = (1 - \beta)(\mathbf{I} - \beta\tilde{\mathbf{S}})^{-1}[\mathbf{A}^{(0)}(\mathbf{s}) - \mathbf{A}^{(0)}(\mathbf{s}')] \quad (74)$$

Since $\|(\mathbf{I} - \beta\tilde{\mathbf{S}})^{-1}\|_2 \leq 1/(1 - \beta\lambda_{\max})$ and vanilla attention is L_0 -Lipschitz:

$$\|\mathbf{A}^*(\mathbf{s}) - \mathbf{A}^*(\mathbf{s}')\|_F \leq \frac{1 - \beta}{1 - \beta\lambda_{\max}} L_0 \|\mathbf{s} - \mathbf{s}'\|_2 \quad (75)$$

$$\leq L \|\mathbf{s} - \mathbf{s}'\|_2 \quad (76)$$

figures/attention_maps.pdf

Figure 3: Attention maps for example mutation (L15V in lysozyme). **(a)** Vanilla attention focuses on sequential neighbors. **(b)** Diffused attention incorporates spatially proximal residues (marked with red boxes) that are distant in sequence. **(c)** 3D structure showing attended residues, confirming spatial proximity.